Adaptive Scheduling in Spark

by

Rohan Mahajan

Submitted to the Department of Electrical Engineering and Computer Science
in partial fulfillment of the requirements for the degree of

Master of Engineering in Computer Science and Engineering

at the Massachusetts Institute of Technology

June 2016

Author . . . . . .
Department of Electrical Engineering and Computer Science
May 20, 2016

Certified by . . . . . .
Prof. Matei Zaharia
Thesis Supervisor

Accepted by . . . . . .
Dr. Christopher J. Terman
Chairman, Masters of Engineering Thesis Committee

Adaptive Scheduling in Spark

by

Rohan Mahajan

Submitted to the Department of Electrical Engineering and Computer Science
on May 20, 2016, in partial fulfillment of the
requirements for the degree of
Master of Engineering in Computer Science and Engineering

Abstract

Because most data processing systems are distributed in nature, data must be transferred between these machines. Currently, Spark, a prominent such system, predetermines the strategies for how this data is to be shuffled but in certain situations, perfomance may be improved by not performing the typical strategy. We add functionality to track metrics about the data during the job and appropriately adapt our shuffle strategy. We show improvements in regular shuffle performance, joins using Spark's RDD interface, and joins in Spark SQl.

Acknowledgments

First, I would like to thank my parents Umesh Mahajan and Manjula Mahajan for their enduring support and love throughout my time at MIT.

I would like to thank Professor Matei Zaharia for his guidance, patience, and support while advising me throughout this project. I learned a lot throughout this project and am extremely grately for the support.

At MIT, my work would never have been completed if not for the support of my friends. I would like to thank them for all the lessons that I have learned and all of the memories that I have created.

Contents

List of Figures

List of Tables

Chapter 1

Introduction

1.1 Spark and MapReduce

New data processing systems such as Spark and MapReduce have been designed to help process the increasing amount of data. Instead of relying on just one powerful computer, these systems use many computers due to lower costs, increased scalability, and improved fault tolerance. Because these systems are distributed in nature, they have stages (shuffle stages) where they transfer information between computers.

1.2 Shuffle

We will use MapReduce to explain the shuffle in more detail, but the main concepts still apply to Spark.

1.2.1 Shuffle Introduction

In the first stage of MapReduce, the map phase, the data is loaded onto different computers and computation is performed on it that results in a group of key-value pairs. The final phase of MapReduce,the reduce phase, assumes that all key-value pairs with the same key are grouped together onto the same machine. We call this property the shuffle guarentee. Thus, the shuffle phase, an intermediate phase that the system handles internally, transfers

key-value pairs between machines to satisfy the shuffle guarentee.

Figure 1-1 display the inner workings of the shuffle phase in MapReduce. For instance, a programmer may want to count the number of letters in a distributed file. The mappers will each load part of the distributed file and count the number of letters in their part. However, the systems needs to aggregate the count for each letter and thus all the counts for letter a will be sent to worker1, letter b will be sent to worker 2, letter c will be sent to worker c. These reducers will then promptly aggregate the counts that they receive from the mappers.

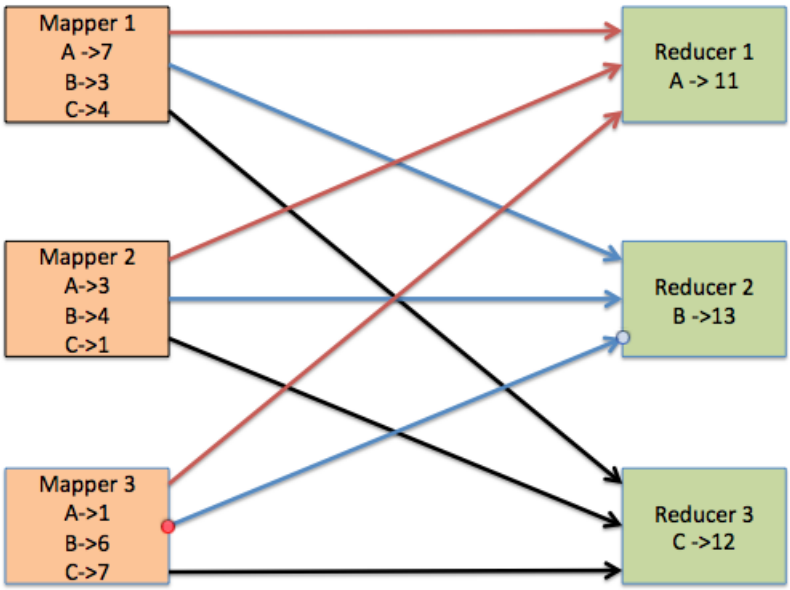

Figure 1-1: Shuffle for Letter Count in MapReduce
This figures demonstrates a basic shuffle in MapReduce. Red arrows indicate the transfer of letter counts of A,blue arrows are used for B, and black arrows are used for C.

Due of the huge amounts of keys, these systems do not transfer data on the granurality of keys. Instead, they use partitions, which contain key-value pairs with different keys. Programmers can pick different partitioning functions such as hash partitiong and range partitiong to map keys to partitions. Two identical keys are guarenteed to be in the same partition. As long as all the mappers partition their data in the same way and send each partition with the same index to the same reducer, the system satisfies the shuffle guarentee.

1.2.2 Shuffle Analysis

MapReduce is constrained by the slowest worker; therefore, minimizing the latency of the slowest worker should improve imperformance. Balancing the amount of data sent to each reducer helps achieve this by reducing both network latency and also the execution time for the slowest worker. Figure **??**, depicts a shuffle scenario that results in unbalanced paritions. Each mapper partition gets sent to the reducer id equal to the partition id mod the number of reducers. This protocol in theory should result in pretty balanced reducers but is not guarenteed to. As depicted, Reducer 3 receives 40MB of data but Reducer 1 receives 90MB of data. However, if we knew the size of each partition after the mappers have run, we could more intelligently balance the reducers. As seen in Figure **??**, with the same map output partitions, the system could attain complete balance of 60MB for each reducer.

1.3 Adaptive Scheduling of Joins

1.3.1 Join Basics

A common operation in these data processing environments is a join. A join basically combines two tables by finding intersections between keys in respective columns. For instance, if we have Table1.1 and Table1.2 that we are trying to join based on the intersection of key1 and key2, the resulting output is Table1.3

Key1	Value1
a	1
a	1
b	3
c	4

Table 1.1: Table for Dataset 1

Key2	Value2
a	5
c	7

Table 1.2: Table for dataset 2

Key1	Value1	Value2
a	1	5
a	2	5
c	4	7

Table 1.3: Table of Joined Data

1.3.2 Shuffle Join

The actual implementation of joins in MapReduce is very similar to the shuffle scenario presented above. Instead of one dataset participating in the shuffle, two datasets participate in the shuffle and ensure that their corresponding partitions are both sent to the same reducer. Figure 1-4 details a shuffle join. For both datasets, all of the keys that mapped to partition 1 were sent to the reducer 1 and this happens respectively for the rest of the partitions.

1.3.3 Broadcast Join

The diagram above may seem to imply that mappers and reducers are different machines. However, this distinction is artificial and there are no seperate machines for mappers and reducers. Therefore, not all data in the shuffle stage is transferred over the network. In Figure 1-1, if Mapper 1 and Reducer 1 were the same machine, the key value pair A=7 would be read locally and not have to be the sent over the network.

Because transferring data over the network could be a bottleneck, the broadcast join tries to increase the amount of data being read locally.For instance, in Figure 1-4, dataset 1 is drastically bigger than dataset 2. As seen in Figure 1-5, the broadcast join keeps the bigger dataset in place and sends the entirety of dataset 2 to every reducer. Even though all of dataset 1 stays in place, this method satisfies theshuffle party because all partitionsof dataset 2 are sent. The diagram shows that the network traffic is reduced from megabytes to kilobytes.

Broadcast Join is not always the optimal strategy. Because the entirety of the smaller dataset is sent to all partitions, the amount of total computation time increases. Additionally, if the datasets are approximately the same size, network traffic will actually increase.

Each join strategy is the optimal strategy in different situations. Thus, it becomes imperative to be able to pick the strategy after the mappers have run and we know the size of the map output files..

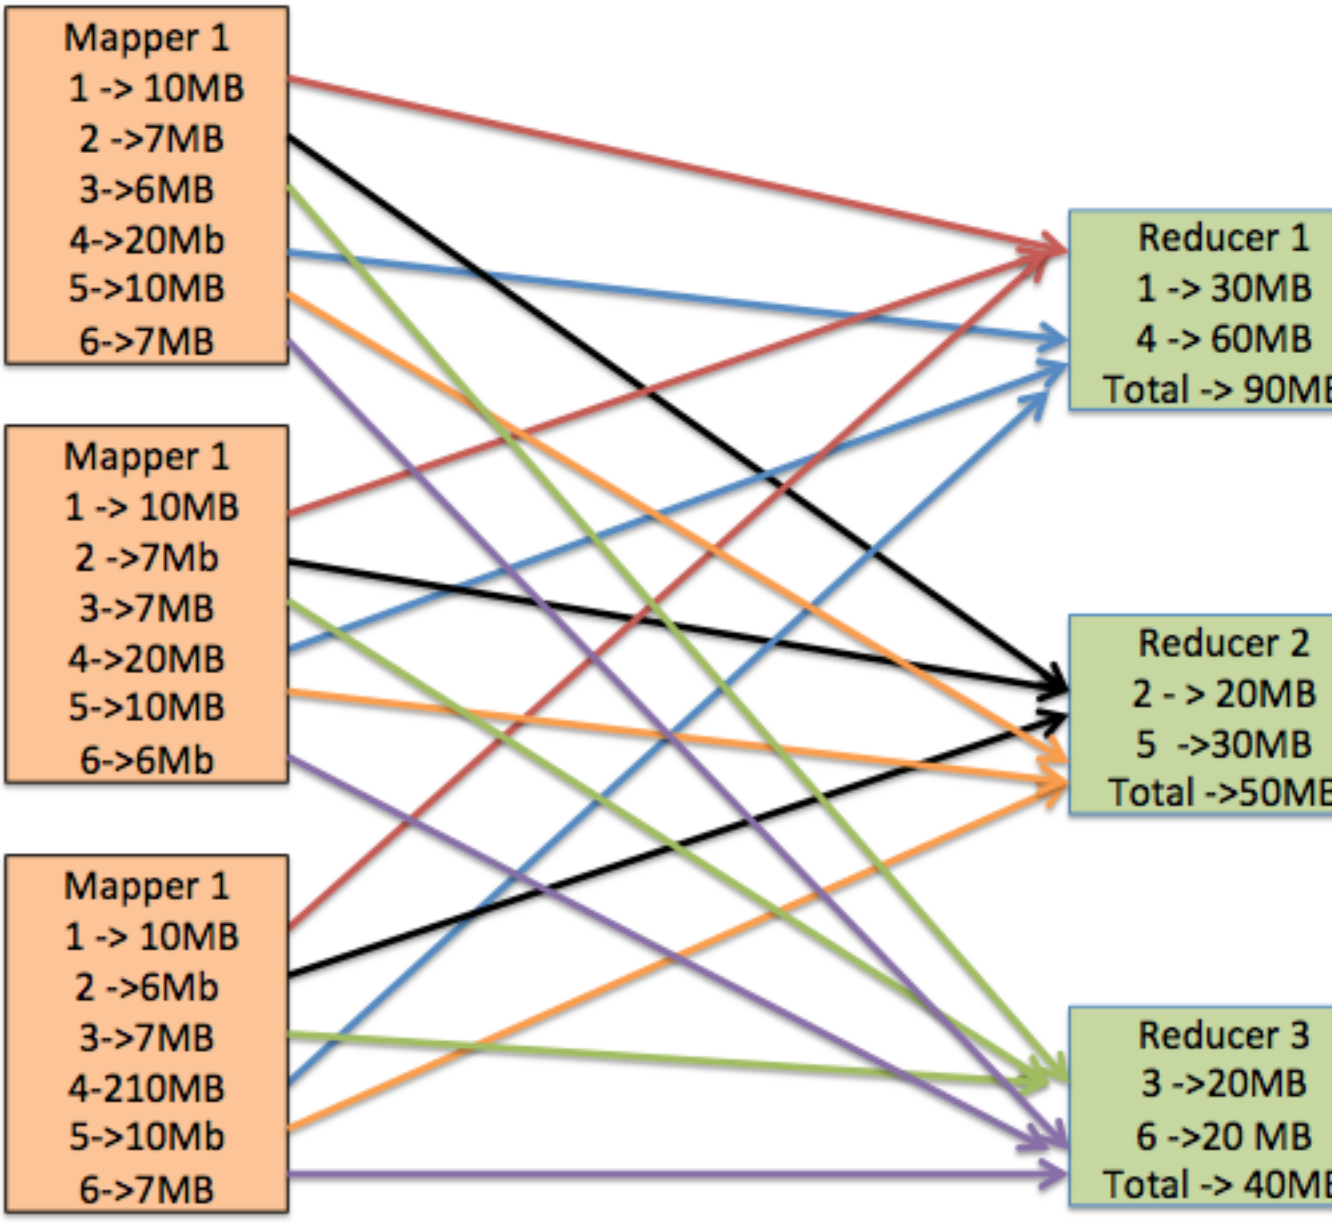

Figure 1-2: Unbalanced shuffle of partitions

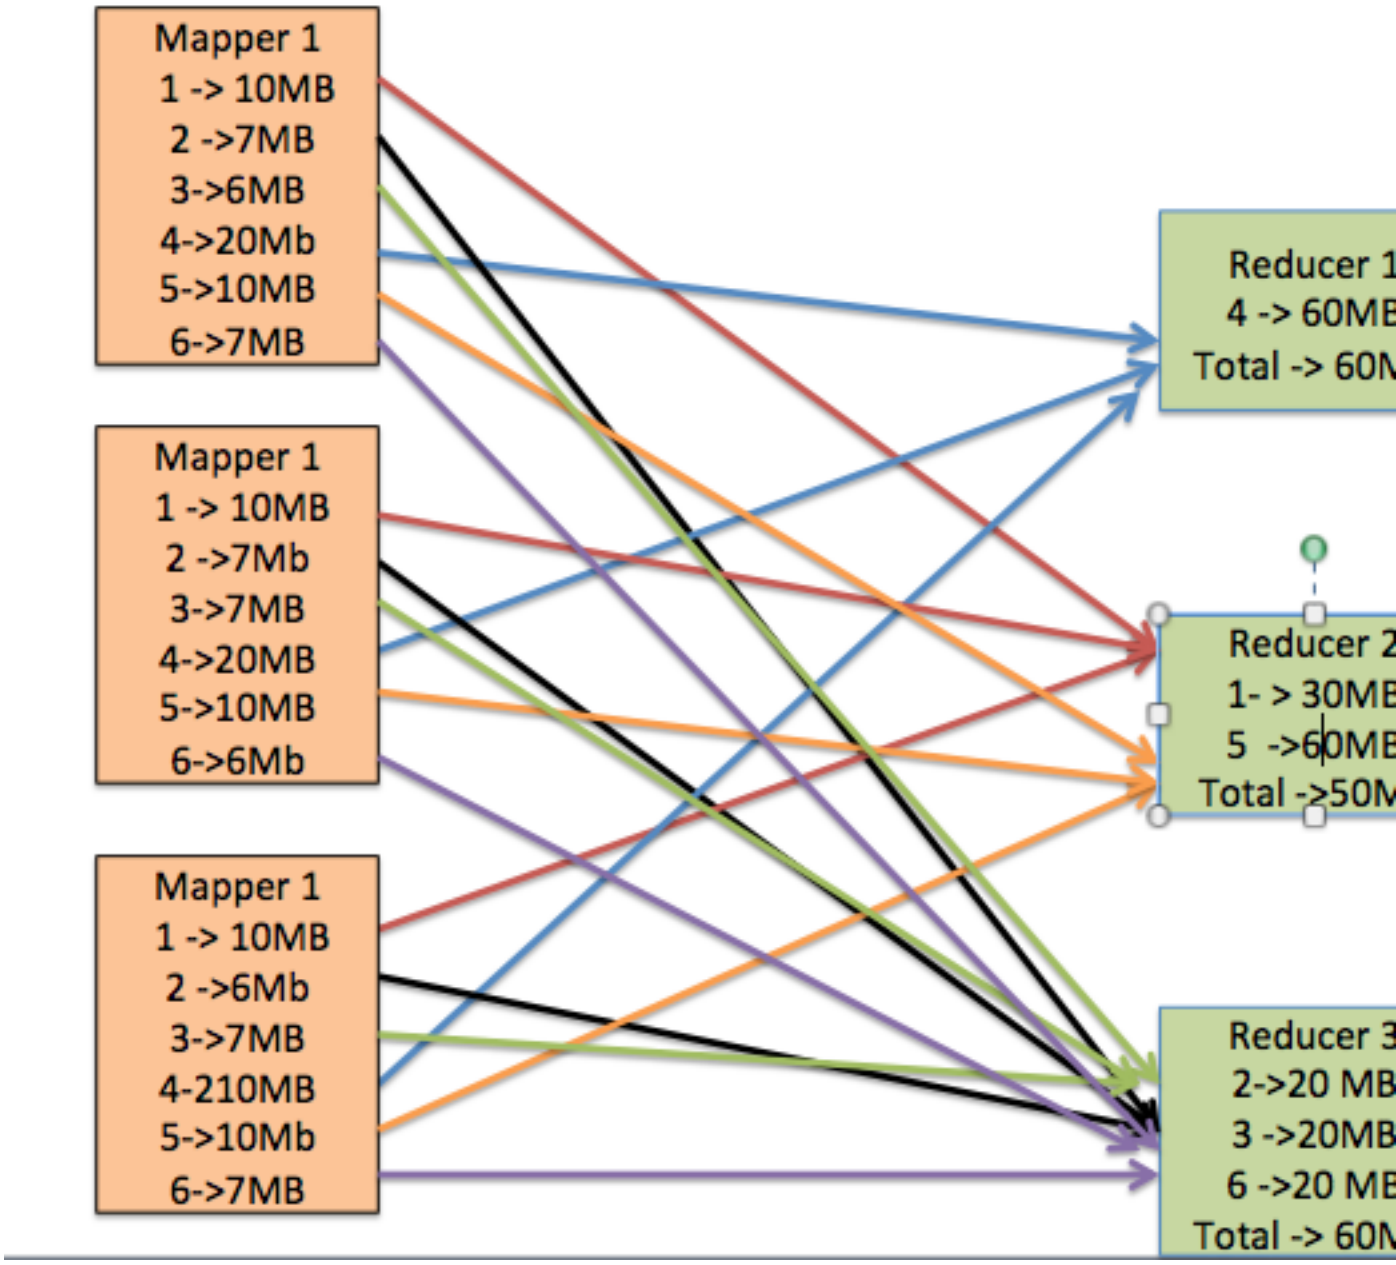

Figure 1-3: Balanced shuffle of partitions.

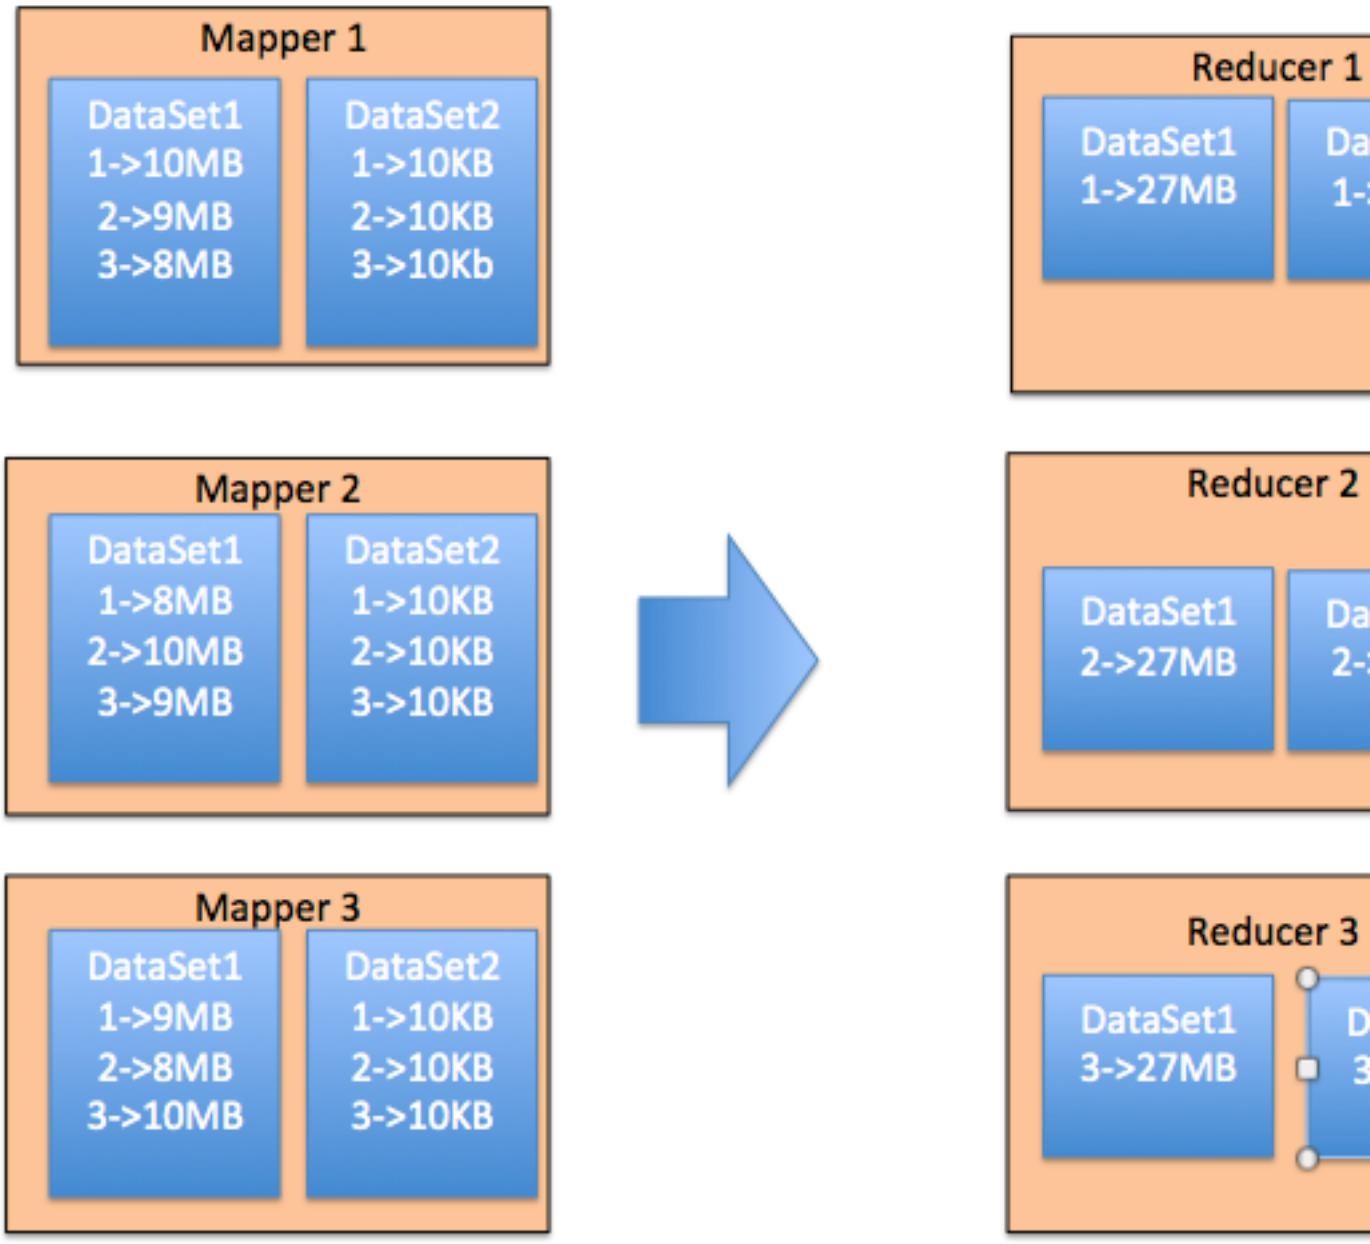

Figure 1-4: Typical Shuffle Join

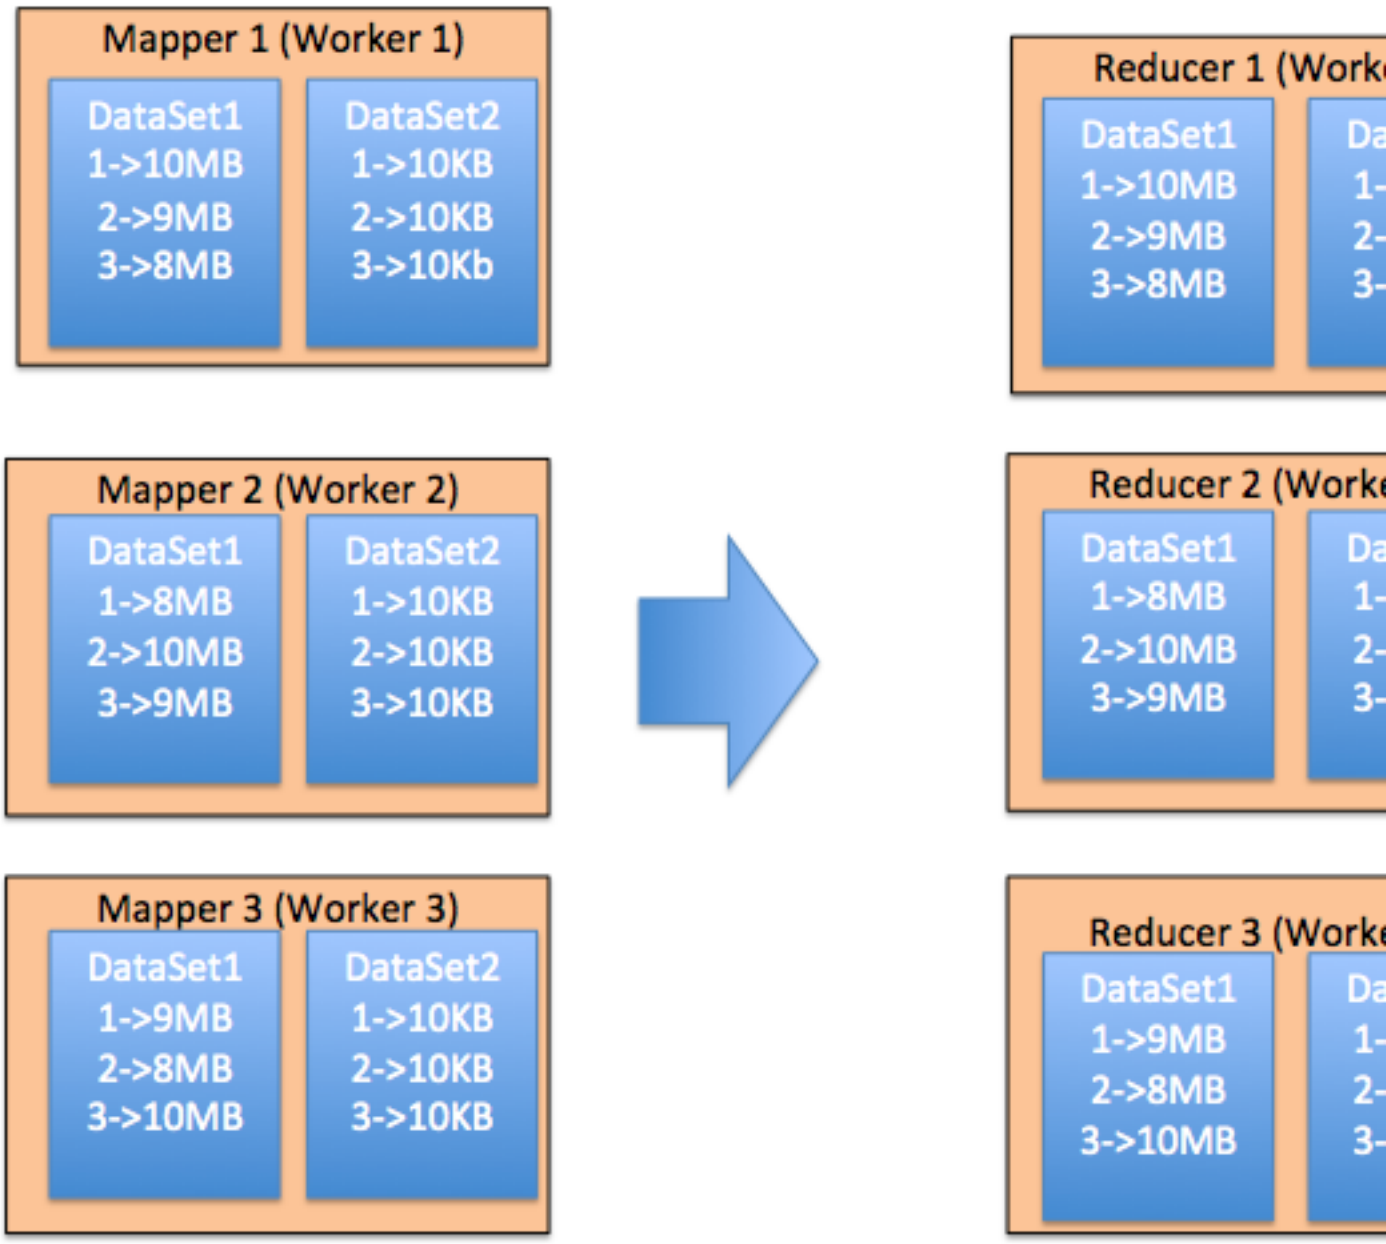

Figure 1-5: Broadcast Join

Chapter 2

Implementation

2.1 Spark

All of the code was implemented in Spark. Although the code was implemented in Spark, it could also be implemented in MapReduce to achieve similar performance improvement. The Resilient Distributed Dataset(RDD) is the main interface within Spark. The RDD can be created from data or from another RDD. The key attributes of an RDD are its inputs, the number of partitions, and how each of its partitions is computed based on its inputs.

Profressor Matei Zaharia added code that allowed the tracking of sizes of map output files.

2.2 ShuffledRDD

The RDD we developed we developed is a new version of ShuffledRDD, ShuffledRDD2. Its inputs are first a shuffle dependency, which is basically a bunch of map output partitions, and second a number of reducers, which indicates the number of partitions for ShuffledRDD@. In The regular ShuffledRDD, each of its partition naively requests a segment of map output partitions as depicted in Figure **??**. ShuffledRDD2 implements the more complicated scheme seen in Figure **??** As this is a proof of concept, each output ShuffledRDD2 partition can only request consecutive map partitions. In other words, it is impossible for a ShuffledRDD2 partition to have map output partitions 1 and 3, without having 2. For this

constraint and the given number of partitions , ShuffleRDD2 is guarenteed to produce the most optimally balanced output.

2.3 Joins

2.3.1 ShuffleReader Changes

As mentioned in the broadcast join section, the bigger RDD must stay in place. The current interface only allows a reducer to request a specific map output partition from all of the mappers. For the bigger RDD, we would thus have to request map output partitions from other machines, which defeats the purpose of the broadcast machine. Thus, we added the capability of requesting a specific partition from just one mapper.

2.3.2 ShuffleJoinRDD and BroadcastJoin RDD

We implement two different type of RDD's, the ShuffleJoinRDD and the BroadcastJoinRDD. Both of these RDD's take two shuffle dependencies, which remember are basically the outputs of map stages, partitioned in a certain way. These dependencies must be partitioned in the same way. Otherwise, we have no way of ensuring that two identical keys are in the same partition.

The ShuffleJoinRDD implementation is very similar to ShuffledRDD. Instead of fetching map output partitions from just one dependency, it fetches the corresponding map output partitions from both dependencies. For instance, ShuffledJoinRDD partition 1 will fetch dataset1 partition 1 and dataset2 partition 1 from all of the workers. Once these partitions are fetched, it create a map with the key value pairs of the smaller partition. IT iterate through the bigger partition, seeing if there are keys present in this map, and if so, we add this to ourput.

The BroadcastJoinRDD implements the broadcast shuffle. For each BroadcasstJoinRDD partition, it requests one local map output partition from the bigger RDD using the

new request capability and all of the paritions from the smaller RDD, thus giving us all of the smaller RDD. We then use the same strategy to actually join the same strategy as the ShuffledJoinRDD to find the intersections.

2.3.3 Joins in Spark SQL

Although the RDD interface is very popular, many programmers and data analysts prefer not to use this interface and are more familiar with the sql and thus Spark offers a sql like interface. One popular operation within sql is join. Although the user still writes in sql, Spark still executes the code using RDD's.

Because we are not just using the RDD interface and Spark automatically converts the sql query into a query plan, the implementation is much more complicated. We only implement our optimization for sort merge join.

Although the exact semantics for how a sort merge join can be found here,the sort merge join requires the shuffle property for the two datasets it is joining. To help achieve this, the sort merge join applies an exchange operator on each of the mapoutputs. These exchange operators produce ShuffleRowRDDS, which for our purposes are equivalent to ShuffledRDDS. In the next stage, each partition in the first ShuffledRowRDD is compared to the partition with the same index in the second ShuffledRowRDD. The only difference between this and how the join RDD's work is pretty semantic in that instead of one RDD requesting partitions from multiple mapper, two RDD's repartition their data and then are compared partition by partion. By default, the code performs a shuffle join almost exactly in a manner with how the ShuffleJoinRDD works. One ShuffleRowRDD requests the corresponding partitions from its mapoutput just like Figure **??** and the other ShuffleRowRDD does the exact same but with its dataset. However, if only one input RDD is smaller then a user configured threshold, we use the broadcast join optimization. The bigger ShuffledRowRDD will be exactly like its parent. The other ShuffledRowRDD will have the same number of partitions as the bigger ShuffledRowRDD with each partition containing the entirety of the smaller input RDD. The correctness guarentees are the same as for join

RDD’s.

Chapter 3

Experiment

3.1 Setup

All jobs were run using the spark/ec2 launch scripts. They were run on four aws m1.large machines. They were run ten times, with the last times being average.

3.2 Regular Shuffle

3.3 Broadcast and ShuffleJoinRDD

3.4 Spark SQL join

Chapter 4

Future Research and Conclusion

4.1 Future Research

4.1.1 Extension of Shuffle

ShuffledRDD2 is limited in a couple ways. First, each reducer can only fetch partitions consecutively, so allowing it to pick non-consecutive partitions could potentially improve performance.Second, the current version only supports inputing the number of reducers. Users could prefer an interface where they input the maximum number of bytes a reducer can have and the system automatically determines the number of reducers.

4.1.2 Extension to Join

First, we implement our changes in the exchange framework to make the easiest possible change to allow for our optmization, but we could conceivably do this in a cleaner manner.

Second, users have to statically pass in thresholds that determine when to switch between broadcast and shuffle joins. The system should automatically determine this based on factors such as the size of the RDDs as well as additional info such as the network bandwith and memory of each machine.

Third, we either broadcast an entire RDD or default to the shuffle pattern. However, if RDD1 has a big partition 1 and a small partition2 and RDD2 has a small partition 1 and big partition 2, the systems performs a shuffle. However, the system could save time by having

RDD1 broadcast its partition 1 and RDD2 broadcast its partition 2.

Fourth, in the broadcast join in Spark SQL, each ShuffleJoinRDD partition requests the entirety of its input. This request is made over the network for each partition, but generally multiple ShuffleJoinRDD partitions are on the same machine. Thus, a request should be made once per machine and stored in memory for the other partitions to use.

4.2 Conclusion

In conclusion, we show that improvements can be made to shuffle stage of Spark. Instead of predetermining our shuffle strategy, we can adapt it based on the output of the mappers. We show that we can use this for improvments in the regular shuffle, in joins with rdds, and in joins using in Spark Sql. Although we have shown improvements, the work can be extended with simple changes to further improve performance.

Bibliography